

Datacern AI

Governed AI for regulated back-office decisions

Datacern lets AI do the regulated work — reading the file, applying the rules, drafting the recommendation with its evidence — while making it **structurally impossible** for the AI to make the final call. Every AI output is a proposal. A named human approves it. That is enforced by the system, not by policy.

HOW IT WORKS

01 Work arrives A claim, dispute or alert flows in from the customer's own systems.	02 AI prepares An agent reads it, applies the rules, drafts a recommendation with citations.	03 A human decides It becomes a proposal. Risky items need a second, different person.	04 Receipt Who proposed, who approved, what changed — permanently recorded.
---	--	--	---

WHAT IT DOES

Casework & triage

Cases anchored to real data rows, SLA clocks, dispositions, evidence attachments and a full timeline.

Approval desk (four-eyes)

Proposals with a system-generated description of the real effect. No self-approval, ever.

Agents

Nine platform agents — triage, analytics, onboarding, dashboards, training, ML engineer, inference, governance, routing — plus no-code tenant agents, each inside an enforced guardrail envelope.

Decision tables

Plain-language *when* → *then* rules over real columns. Deterministic, instant, no LLM, and the fired rule is named.

Data platform

18+ source connectors, healthcare/finance standards (X12, HL7/FHIR), versioned lakehouse datasets, lineage, a governed semantic layer, dashboards and entity resolution.

Machine learning

21 algorithm templates, visual pipelines, experiment tracking, batch scoring, and model promotion behind a human approval gate.

The learning loop

Every human correction becomes a training example, distilled into smaller models — so decisions get cheaper and better over time.

Governance & audit

Tamper-evident, WORM-exported audit trail; tenant isolation enforced by the database; per-request policy checks; cost, budget and ROI reporting.

Capability packs

28 pre-built kits — card disputes, AML, claims, prior auth, invoice audit and more. A pack refuses to install until the customer's real data matches what the workflow needs.

Deployment

Runs on AWS, Azure, GCP or your own cluster. Bring your own identity, secrets and observability.

WHO USES IT

The specialist

Opens a queue with a drafted recommendation on every item. Accepts, edits or escalates.

The manager

Works the approval desk — approves or rejects with a reason, and cannot approve their own.

The compliance lead

Never touches a case. Pulls the trail when an examiner asks how a decision was made.

23

services

28

industry packs

9

AI agents

21

ML algorithms

18+

data connectors

100%

decisions
human-approved

WHERE WE ARE

Pre-launch and candid about it. The platform runs end-to-end on real infrastructure — real logins, real data, a real language model, real tenant isolation — and has been watched doing so. It has not yet run outside a development machine: no customers, no load testing, no independent security audit. Unattended AI runs are limited; human-initiated work is the solid path. We would rather show you the working system and name the gaps than pitch a roadmap.